

Visualizing Spatial Data

Angela Li, Center for Spatial Data Science

January 21, 2020

Outline

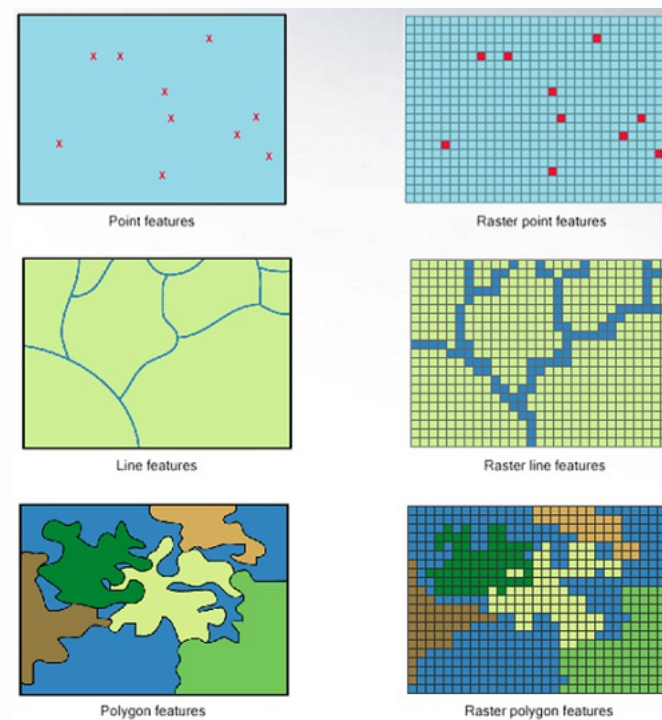
1. What Is Spatial Data?
2. Reading/Writing Spatial Data
3. Spatial Data Visualization
4. Spatial Data Transformation / Geoprocessing (if time)

What is Spatial Data?

Vector vs. Raster Data

Spatial data comes in two main formats: **vector** and **raster** data.

Vector data often represents discrete objects, while raster data often represents continuous surfaces.



Vector Data: Geocoding

Sometimes people give me a list of addresses and ask me to map it. I can't do that unless I have the **latitude and longitude**, so I'll need to **geocode** the addresses. For states (polygons), I'll need to find a **geospatial boundary file** with that information (Google around!)

```
bikeshare_addresses <- read.csv("data/bike_addresses.csv")  
head(bikeshare_addresses)
```

##	ID	ADDRESS
## 1	550	McLean Metro
## 2	551	Trinidad Rec Center
## 3	552	Rosedale Rec Center
## 4	553	11th & C St SE
## 5	554	New Hampshire & Gallatin St NW
## 6	555	United Medical Center

Geocoding in R

There are a few options, which usually limit you to 2000 queries. Here's one nice one that uses [OpenStreetMap](#).

```
tmaptools::geocode_OSM("McLean Metro")
```

```
## $query
## [1] "McLean Metro"
##
## $coords
##           x           y
## -77.20791  38.92379
##
## $bbox
##      xmin      ymin      xmax      ymax
## -77.20796  38.92374 -77.20786  38.92384
```

Geocoding will give me latitude and longitude for points

```
bikeshare_latlon <- read.csv("data/bike_addresses_latlon.csv")  
head(bikeshare_latlon)
```

```
##      ID                ADDRESS LATITUDE LONGITUDE  
## 1  550             McLean Metro 38.92400 -77.20813  
## 2  551      Trinidad Rec Center 38.90630 -76.98322  
## 3  552      Rosedale Rec Center 38.89781 -76.97963  
## 4  553             11th & C St SE 38.88591 -76.99148  
## 5  554 New Hampshire & Gallatin St NW 38.95160 -77.01281  
## 6  555      United Medical Center 38.83574 -76.98314  
##                                geometry  
## 1  c(-77.2081293496828, 38.9240097850969)  
## 2      c(-76.983223279, 38.9063067901138)  
## 3  c(-76.9796362771458, 38.8978157882255)  
## 4   c(-76.9914782802471, 38.885915785567)  
## 5   c(-77.012810290596, 38.9516037980159)  
## 6  c(-76.9831462746571, 38.8357447756218)
```

Read in spatial data

To read in spatial data, use the `st_read()` or `read_sf()` function from the `sf` package:

```
library(sf)
```

```
## Linking to GEOS 3.7.2, GDAL 2.4.2, PROJ 5.2.0
```

```
st_read("data/dc_wards.shp")
```

```
## Reading layer `dc_wards' from data source `/Users/angela/Desktop/R-Projects/aaas-  
## Simple feature collection with 8 features and 82 fields  
## geometry type: POLYGON  
## dimension: XY  
## bbox: xmin: -77.1198 ymin: 38.79164 xmax: -76.90915 ymax: 38.99597  
## epsg (SRID): 4326  
## proj4string: +proj=longlat +datum=WGS84 +no_defs
```

```
read_sf("data/dc_wards.shp") # if you like tidy data
```

```
## Simple feature collection with 8 features and 82 fields  
## geometry type: POLYGON  
## dimension: XY  
## bbox: xmin: -77.1198 ymin: 38.79164 xmax: -76.90915 ymax: 38.99597  
## epsg (SRID): 4326
```


Writing spatial data

To write spatial data, use `st_write()` or `write_sf()`:

```
st_write(dc_wards, "data-output/dc_wards.shp")  
write_sf(dc_wards, "data-output/dc_wards.shp")
```


Spatial Data Visualization

Spatial Data Visualization

- Base plotting
- tmap
- ggplot2
- mapview

Base Plotting

```
plot(dc_wards)
```



Plot just the geometry


```
plot(st_geometry(dc_wards))
```



Plot more than one layer


```
bikeshare_locations <- read_sf("data/bikeshare_locations.shp")
```

```
plot(st_geometry(dc_wards))  
plot(st_geometry(bikeshare_locations), add = TRUE)
```



Change the color

```
plot(st_geometry(dc_wards))  
plot(st_geometry(bikeshare_locations), add = TRUE, col = "red")
```



Common issues: Projections and Coordinate Reference Systems

If your map layers won't plot on top of each other, you need to check that they are in the same map projection.

```
plot(st_geometry(dc_wards_proj))  
plot(st_geometry(bikeshare_locations), add = TRUE)
```



Check the projection

These aren't the same!

```
st_crs(dc_wards_proj)
```

```
## Coordinate Reference System:
```

```
##   EPSG: 6654
```

```
##   proj4string: "+proj=utm +zone=11 +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +units=m +v
```

```
st_crs(bikeshare_locations)
```

```
## Coordinate Reference System:
```

```
##   EPSG: 4326
```


```
##   proj4string: "+proj=longlat +datum=WGS84 +no_defs"
```

Project the data

```
bikeshare_locations <- st_transform(bikeshare_locations, 6654)
```

Other R map package options

```
library(tmap)  
qtm(dc_wards)
```




```
tm_shape(dc_wards) +  
  tm_polygons()
```



```
tm_shape(dc_wards) +  
  tm_borders()
```



```
tm_shape(dc_wards) +  
  tm_fill()
```




Change to an interactive mode

```
tmap_mode("view")
```

```
## tmap mode set to interactive viewing
```


```
tm_shape(dc_wards) +  
  tm_polygons()
```



```
tmap_mode("plot")
```

```
## tmap mode set to plotting
```

```
tm_shape(dc_wards) +  
  tm_polygons() +  
  tm_shape(bikeshare_locations) +  
  tm_dots(size = 0.1)
```




More complicated maps possible

```
tmap_mode("plot")
```

```
## tmap mode set to plotting
```

```
tm_shape(dc_wards) +  
  tm_polygons("POP_2000") +  
  tm_shape(bikeshare_locations) +  
  tm_dots(size = 0.1)
```




ggplot2 is also an option

More on this at "[Drawing beautiful maps programmatically with R, sf and ggplot2](#)".

```
library(ggplot2)
```

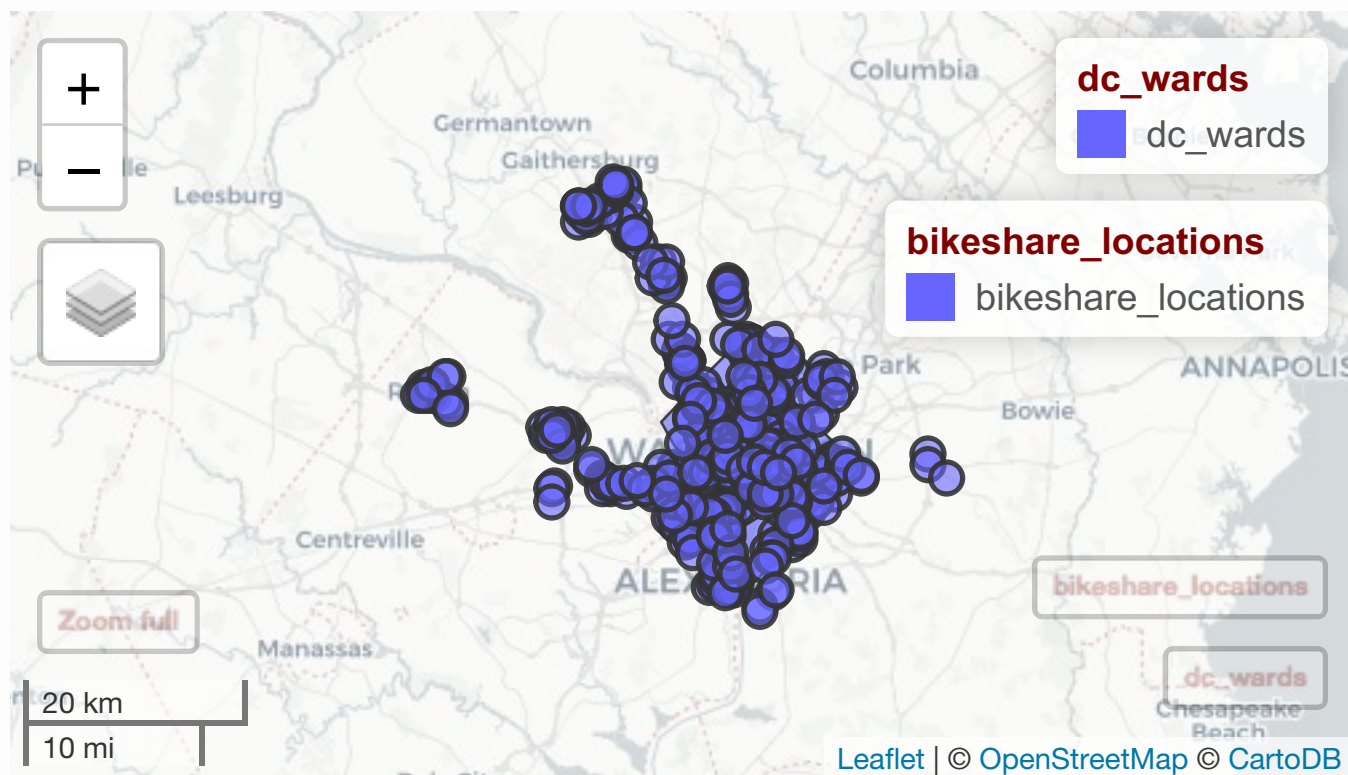
```
## Registered S3 methods overwritten by 'ggplot2':  
##   method      from  
##   [.quosures  rlang  
##   c.quosures  rlang  
##   print.quosures rlang
```

```
ggplot(data = dc_wards) +  
  geom_sf()
```



For an interactive map experience similar to a GIS, check out mapview

```
library(mapview)  
mapview(dc_wards) +  
  mapview(bikeshare_locations)
```



Spatial Data Transformation / Geoprocessing


1. Crop locations to DC Wards
2. Count how many in each ward
3. Calculate bikeshare density by population

Crop locations to DC Wards

```
bikes_in_dc <- st_intersection(bikeshare_locations, dc_wards)
```

```
## although coordinates are longitude/latitude, st_intersection assumes that they are
```

```
plot(st_geometry(dc_wards))  
plot(st_geometry(bikes_in_dc), col = "red", add = TRUE)
```



Count up number of bikes in wards

Two ways to do this...

1. Attribute join

```
st_intersection(dc_wards, bikeshare_locations) %>%  
  dplyr::count(WARD) %>%  
  st_drop_geometry() %>%  
  dplyr::right_join(dc_wards)
```

although coordinates are longitude/latitude, st_intersection assumes that they are

Joining, by = "WARD"

A tibble: 8 x 84

```
##   WARD      n OBJECTID NAME  REP_NAME WEB_URL  REP_PHONE REP_EMAIL  
##   <dbl> <int>   <dbl> <chr> <chr>    <chr>    <chr>    <chr>  
## 1     8    22        1 Ward... Trayon ... http:/... (202) 72... twhite@d...  
## 2     6    68        2 Ward... Charles... http:/... (202) 72... callen@d...  
## 3     7    19        3 Ward... Vincent... http:/... (202) 72... vgray@dc...  
## 4     2    93        4 Ward... Jack Ev... http:/... (202) 72... jevans@d...  
## 5     1    33        5 Ward... Brianne... http:/... (202) 72... bnadeau@...  
## 6     5    32        6 Ward... Kenyan ... http:/... (202) 72... kmcduffi...  
## 7     3    21        7 Ward... Mary M.... http:/... (202) 72... mcheh@dc...  
## 8     4    18        8 Ward... Brandon... http:/... (202) 72... btodd@dc...
```

```
## # A tibble: 8 x 84  
##   WARD      n OBJECTID NAME  REP_NAME WEB_URL  REP_PHONE REP_EMAIL  
##   <dbl> <int>   <dbl> <chr> <chr>    <chr>    <chr>    <chr>  
## 1     8    22        1 Ward... Trayon ... http:/... (202) 72... twhite@d...  
## 2     6    68        2 Ward... Charles... http:/... (202) 72... callen@d...  
## 3     7    19        3 Ward... Vincent... http:/... (202) 72... vgray@dc...  
## 4     2    93        4 Ward... Jack Ev... http:/... (202) 72... jevans@d...  
## 5     1    33        5 Ward... Brianne... http:/... (202) 72... bnadeau@...  
## 6     5    32        6 Ward... Kenyan ... http:/... (202) 72... kmcduffi...  
## 7     3    21        7 Ward... Mary M.... http:/... (202) 72... mcheh@dc...  
## 8     4    18        8 Ward... Brandon... http:/... (202) 72... btodd@dc...
```

Count up number of bikes in wards

Less steps -

1. Spatial join (use the geometry to perform a join):

```
st_join(dc_wards, bikeshare_locations) %>%  
  dplyr::count(WARD)
```

```
## although coordinates are longitude/latitude, st_intersects assumes that they are
```

```
## Simple feature collection with 8 features and 2 fields
```

```
## geometry type: POLYGON
```

```
## dimension: XY
```

```
## bbox: xmin: -77.1198 ymin: 38.79164 xmax: -76.90915 ymax: 38.99597
```

```
## epsg (SRID): 4326
```

```
## proj4string: +proj=longlat +datum=WGS84 +no_defs
```

```
## # A tibble: 8 x 3
```

```
##   WARD      n geometry  
## * <dbl> <int> <POLYGON [°]>
```

```
## 1      1     33 ((-77.03523 38.93743, -77.0348 38.93743, -77.03436 38.93743,...
```

```
## 2      2     93 ((-77.04946 38.91999, -77.04919 38.91954, -77.04918 38.91952...
```

```
## 3      3     21 ((-77.05808 38.95676, -77.05807 38.95672, -77.05805 38.95672...
```

```
## 4      4     18 ((-77.04097 38.99597, -76.99144 38.9573, -76.99163 38.95726,...
```

```
## 5      5     32 ((-76.99144 38.9573, -76.94186 38.91854, -76.942 38.91842, -...
```

```
## 6      6     68 ((-77.0179 38.9141, -77.01786 38.914, -77.01784 38.91393, -7...
```

Find bike density in each ward

```
st_join(dc_wards, bikeshare_locations) %>%  
  dplyr::count(WARD, POP_2011_2)
```

although coordinates are longitude/latitude, st_intersects assumes that they are

Simple feature collection with 8 features and 3 fields

geometry type: POLYGON

dimension: XY

bbox: xmin: -77.1198 ymin: 38.79164 xmax: -76.90915 ymax: 38.99597

epsg (SRID): 4326

proj4string: +proj=longlat +datum=WGS84 +no_defs

A tibble: 8 x 4

##	WARD	POP_2011_2	n	geometry
##	* <dbl>	<dbl>	<int>	<POLYGON [°]>
##	1	1	82859	33 ((-77.03523 38.93743, -77.0348 38.93743, -77.0343...
##	2	2	77645	93 ((-77.04946 38.91999, -77.04919 38.91954, -77.049...
##	3	3	83152	21 ((-77.05808 38.95676, -77.05807 38.95672, -77.058...
##	4	4	83066	18 ((-77.04097 38.99597, -76.99144 38.9573, -76.9916...
##	5	5	82049	32 ((-76.99144 38.9573, -76.94186 38.91854, -76.942 ...
##	6	6	84290	68 ((-77.0179 38.9141, -77.01786 38.914, -77.01784 3...
##	7	7	73290	19 ((-76.94186 38.91854, -76.90915 38.89293, -76.961...
##	8	8	81133	22 ((-76.97229 38.87286, -76.97223 38.87273, -76.972...


```
st_join(dc_wards, bikeshare_locations) %>%
```


Make a map!

```
bikeshare_density_per_ward <-  
st_join(dc_wards, bikeshare_locations) %>%  
  dplyr::count(WARD, POP_2011_2) %>%  
  dplyr::mutate(bikeshare_density = n / POP_2011_2)
```

although coordinates are longitude/latitude, st_intersects assumes that they are

```
tm_shape(bikeshare_density_per_ward) +  
  tm_polygons("n")
```



Resources

- [Tutorials developed by my research center](#)
- [Geocomputation with R](#)
- [Spatial Data Science](#)
- [Data Carpentry Geospatial Lesson](#) (focuses on raster data)